

REPRODUCIBILITY IN COMPUTATIONAL RESEARCH

Logan Ward
Asst. Computational Scientist
Argonne National Laboratory

15 January 2022

Computational Research *Should* Be Always Reproducible

Why? Every little bit is broken down into precise instructions, executed by something else

Why not?

- Some algorithms are random: different results each time
- Computers do not record every bit of data they process

What is your role? Record what (for computer) you did and why (for humans)

ITEM 1: RECORD YOUR WORK IN NOTEBOOKS!

What is Jupyter? Why care?

- Jupyter is environment designed for reproducible computational science
- Stores code with outputs and documentation in a single notebook
 - Modeled after Mathematica
 - <https://www.theatlantic.com/science/archive/2018/04/the-scientific-paper-is-obsolete/556676/>
- Why do I care? Jupyter lets me...
 - organize my code
 - easily run on a remote system
 - keep track of results and rationale
 - communicate my research better



Generate Representations

In this notebook, we generate the inputs for the machine learning model by computing the features for all available data, and generated the pickle file for the data.

```
In [1]: %matplotlib inline
import sys
from glob import glob
sys.path = glob('/projects/CSC249ADCD01/packages/*') + sys.path
from matplotlib import pyplot as plt
import os
import pandas as pd
import numpy as np
import pickle as pkl
import gzip
from matminer.featurizers.base import MultipleFeaturizer
from stopping_power_ml.io import get_charge_density_integrated
from stopping_power_ml.features import LocalChargeDensity
from tqdm import tqdm_notebook as tqdm
```

Configure Parsl

```
In [2]: import parsl
from parsl import *
workers = IPyParallelExecutor()
dfk = DataFlowKernel(executors=[workers])
```

/home/wardlt/miniconda3/lib/python3.6/site-packages/ipykernel

Controller appears to be listening on localhost. If this is true, you should specify Client(d.org') or instruct your controller to listen on an

How I write a notebook

One notebook per "experiment" or "idea"

1. Title and short abstract
"What am I doing here and why"
2. Load in libraries
Put them up front, so it crashes early
3. Load in data from disk
Ex: training data, results from other notebook
4. Each step in their own block
 - Introduction
 - Code and explanation
 - Figure/visualization
 - Explanation of finding

Example Step

Introduction

Simple Test: Channel

Here, we have a particle traveling forward or backwards along the channel of FCC AI. The path forward and backwards are identical, so we should get the same stopping power

```
In [5]: @App('python', dfk)
def compute_stopping_power(starting_point, direction, traj_computer=traj_computer):
    return traj_computer.compute_stopping_power(starting_point, direction, 1)
```

Set up calculations for the stopping power in the channel

```
In [6]: forward = compute_stopping_power([0,0.75,0.75], [1,0,0])
```

```
In [7]: backward = compute_stopping_power([0,0.75,0.75], [-1,0,0])
```

Wait for them to finish

Explanation

```
In [8]: %%time
forward = forward.result(); backward = backward.result()
```

CPU times: user 12 ms, sys: 12 ms, total: 24 ms
Wall time: 1.8 s

```
In [9]: print('Forward stopping power: ', forward[0])
print('Backward stopping power: ', backward[0])
print('Difference: ', backward[0]-forward[0])
```

Forward stopping power: 0.2343000208236084
Backward stopping power: 0.23431760358559353
Difference: 1.758276198512987e-05

Visualization and conclusion

Finding: Consistent with my initial expectations, they are indeed the same (within numerical tolerances).

Someone should be able to understand this without knowing Python!

The Full Narrative

Assessing Forward/Backward Asymmetry in Stopping Powers

Our stopping power model predicts different stopping powers for particles traveling forward than those traveling backwards on some trajectories. The purpose of this notebook is to showcase why this is valid.

```
In [1]: %matplotlib inline
from matplotlib import pyplot as plt
from stopping_power_ml.features import LocalChargeDensity
import pandas as pd
import numpy as np
import pickle as pkl
import os
```

Configure parsl

```
In [2]: import parsl
from parsl import *
#from parsl_config import config
workers = IPyParallelExecutor()
dfk = DataFlowKernel(executors=[workers])
print("Parsl version : ", parsl.__version__)
```

Parsl version : 0.5.0

/home/wardlt/miniconda3/lib/python3.6/site-packages/ipyparallel/client/client.py:458: RuntimeWarning:

Controller appears to be listening on localhost, but not on this machine.
If this is true, you should specify Client(...,sshserver='you@js-168-224.jetstream-cloud.org')
or instruct your controller to listen on an external IP.

Load in the Tools

We'll need the trajectory computer, and the charge density so that we can make illustrative plots.

```
traj_computer = pkl.load(open('traj_computer.pkl', 'rb'))
```

```
charge_density = pkl.load(open(os.path.join '..', 'density_interp.pkl', 'rb'))
```

[Link to GitHub Page](#)

Pitfalls with Notebooks

1. Not including documentation

Advice: Write what you're going to do first

2. Notebook not capturing entire process

Problem: Jupyter lets you execute cells out of order

Advice: Periodically "Restart Kernel and Run All Cells"

3. Duplicate code between notebooks

Advice: Make a separate module for common code

4. Library conflicts

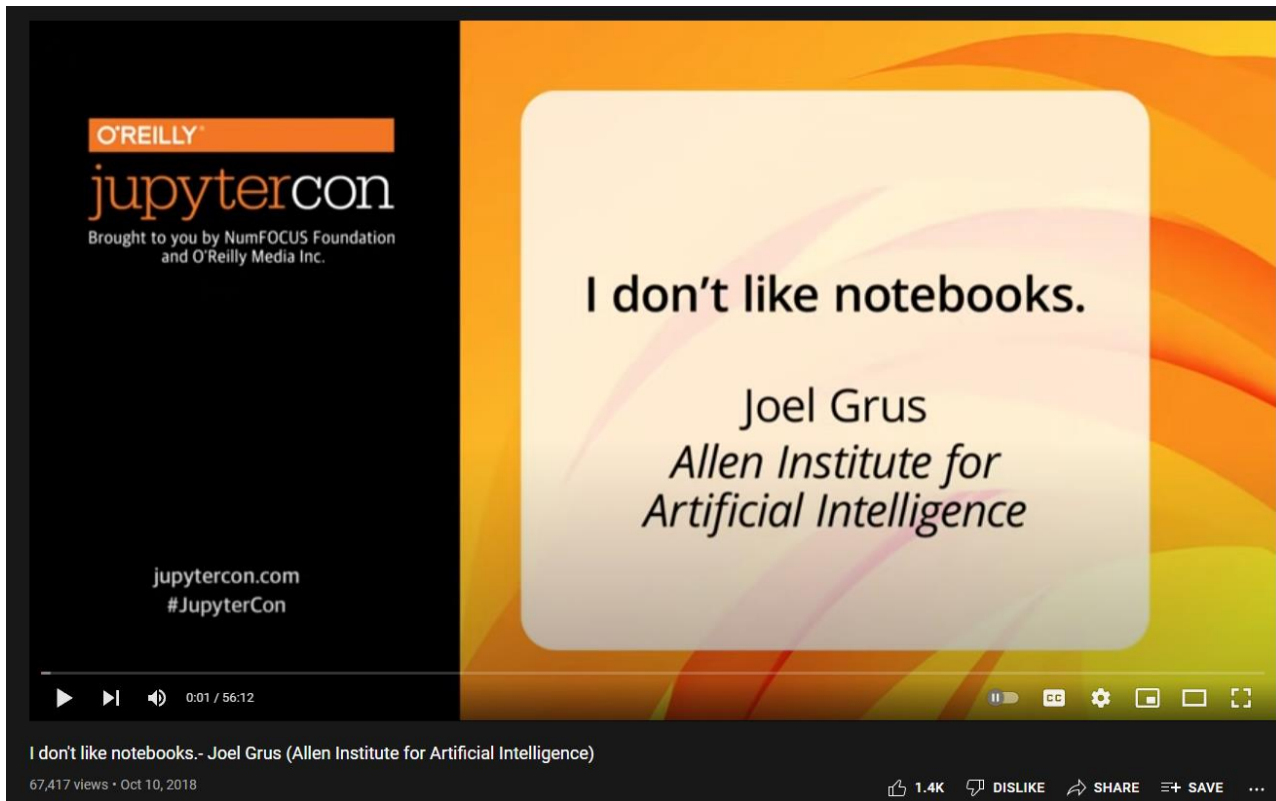
Advice: Run mature projects in container, separate machines

Advice: Make an "environment.yml" file

5. The "one cell notebooks"

Advice: <10 lines of code per cell

More Pitfalls with Notebooks



The image shows a YouTube video player interface. The video content is a presentation slide with an orange and yellow wavy background. The slide text reads: "I don't like notebooks." followed by "Joel Grus" and "Allen Institute for Artificial Intelligence" in italics. The left side of the video player has a black overlay with the O'Reilly logo, "jupytercon" text, and "Brought to you by NumFOCUS Foundation and O'Reilly Media Inc." Below this, it says "jupytercon.com" and "#JupyterCon". The video player controls at the bottom show a progress bar at 0:01 / 56:12, a play button, and various icons for settings, captions, and full screen. The video title "I don't like notebooks.- Joel Grus (Allen Institute for Artificial Intelligence)" is displayed below the player, along with "67,417 views • Oct 10, 2018". At the bottom right, there are icons for likes (1.4K), dislikes, share, save, and a menu icon.

O'REILLY
jupytercon
Brought to you by NumFOCUS Foundation
and O'Reilly Media Inc.

jupytercon.com
#JupyterCon

I don't like notebooks.

Joel Grus
*Allen Institute for
Artificial Intelligence*

▶ ⏮ 🔊 0:01 / 56:12

⏸ CC ⚙️ 📺 🖥️ 🗑️

I don't like notebooks.- Joel Grus (Allen Institute for Artificial Intelligence)

67,417 views • Oct 10, 2018

👍 1.4K 🗑️ DISLIKE ➦ SHARE ➦ SAVE ...

Organizing Multiple Notebooks

🏠 > single-velocity

Name	Last Modified
convergence-detection	a month ago
data	a month ago
feature-analysis	a month ago
figures	a month ago
manifold	17 hours ago
neural-network	15 hours ago
runinfo	15 hours ago
0_collect-subset.ipynb	16 hours ago
1_build-machine-learn...	16 hours ago
2_evaluate-direction-d...	16 hours ago
3_forward-backward-as...	in a few seconds
best_model.pkl	16 hours ago
best_weight.pkl	a month ago
direction_stopping.pkl	16 hours ago
run-notebooks.sh	a month ago
stopping_power_result...	16 hours ago
traj_computer.pkl	16 hours ago

Re-usable datasets in special folder

Figures in special folder (with meaningful names!)
(make pub-ready figures!)

Dependent tasks in subfolders

Notebooks numbered by execution order

Results shared between steps using pickled files

A bash script to re-run all notebooks

Use Github as a Lab Notebooks

The screenshot displays a GitHub repository interface. At the top, the commit history is visible, listing three commits by user WardLT: 'Switched to using spherical coordinates for traj optimizer' (3 days ago), 'Use global optimizer to find best trajectory' (4 days ago), and 'Added initial attempt to get "optimized" trajectories' (7 days ago). A red-bordered box highlights the first commit with the text 'Use GitHub to track changes'. Below the commit history, the file path 'stopping-power-ml / 1_compute-representation.ipynb' is shown. A commit message 'WardLT Switched back to using time offsets for the change density' is visible, along with the commit hash '6d7d3f2' and the time '7 days ago'. A red-bordered box highlights this section with the text 'Make branches for dead-ends'. The file statistics show '459 lines (458 sloc)' and '86.5 KB'. A red-bordered box at the bottom right contains the text 'GitHub renders your notebooks: Great for Sharing with Collaborators!'. At the bottom left, the text 'Generate Representations' is displayed. A central text overlay reads 'Keep code, notes, figures, timelines in one spot'.

Commits on May 21, 2018

Switched to using spherical coordinates for traj optimizer
WardLT committed 3 days ago

Commits on May 20, 2018

Use global optimizer to find best trajectory ...
WardLT committed 4 days ago

Commits on May 17, 2018

Added initial attempt to get "optimized" trajectories
WardLT committed 7 days ago

Branch: master ▾ stopping-power-ml / 1_compute-representation.ipynb

Find file Copy path

WardLT Switched back to using time offsets for the change density 6d7d3f2 7 days ago

1 contributor

459 lines (458 sloc) | 86.5 KB

Generate Representations

Use GitHub to track changes

Make branches for dead-ends

Keep code, notes, figures, timelines in one spot

GitHub renders your notebooks:
Great for Sharing with Collaborators!

Takeaway for Using Jupyter

1. **Write your notebooks like papers**
Explain what you are doing, why, and what you found
Learn Markdown to include links, equations, pictures
2. **Break up complex projects into multiple steps**
Each notebook should tell a single story
3. **Notebooks should be easy to re-run**
Use "Restart and Run," bash scripts
4. **Track your changes with Git**
Explain what changes you made, and why

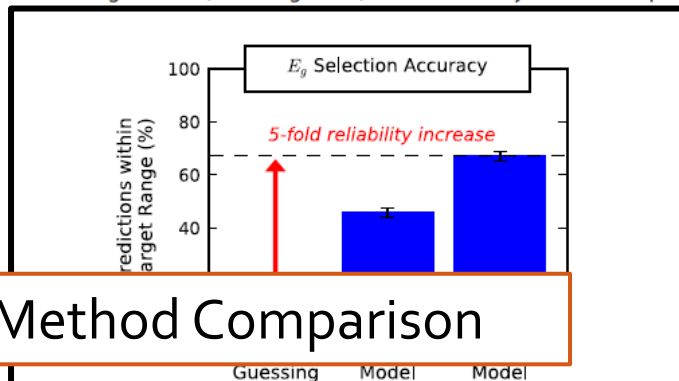
ITEM 2: PUBLICATION IS NOT JUST FOR PAPERS!

Is my 2016 paper reproducible?

ARTICLE OPEN

A general-purpose machine learning framework for predicting properties of inorganic materials

Logan Ward¹, Ankit Agrawal², Alok Choudhary² and Christopher Wolverton¹



Method Comparison

Figure 1. Performance of three different strategies to locate compounds with a band gap energy within a desired range: randomly selecting nonmetal-containing compounds, and two strategies using the machine-learning-based method presented in this work. The first machine learning strategy used a single model trained on the computed band gap energies of 22,667 compounds

that use machine learning to predict properties of materials. This work demonstrated successful prediction of band gap energies for a wide range of materials. To enable

Table 2. Compositions and predicted band gap energies of materials predicted using machine learning to be candidates for solar cell applications

Composition	E_g (eV)
ScHg ₄ Cl ₇	1.26
	1.16
	1.28
	1.11
	1.19

Predicted Materials

Abbreviations: DFT, density functional theory; OQMD, open quantum materials database.

Compositions represent the nominal compositions of novel ternary compounds predicted by using methods developed in ref. 15. Band gap energies were predicted using a machine learning model trained on DFT band gap energies from the OQMD² using methods described in this work.

I released the code, but not *ALL* of it

(Spoiler: I regret it!)

From Original Source Code

Supplementary information


 PDF files

1. [Supplementary Information](#)

 Zip files

1. [Supplementary Information](#)

... to New Notebooks

 jupyter

Files

Running

Clusters

Select items to perform actions on them.


☐ 0  / data / [predicting-band-gap-energies](#)

 ..

☐  [build-and-test-hierarchical-model.ipynb](#)

☐  [evaluate-meredig-predictions.ipynb](#)

☐  [identify-solar-cell-materials.ipynb](#)

☐  [plot-cv-results.ipynb](#)

☐  [bandgap-model-dataset-template.obl](#)

... and there are differences!

Reported Results (2016)

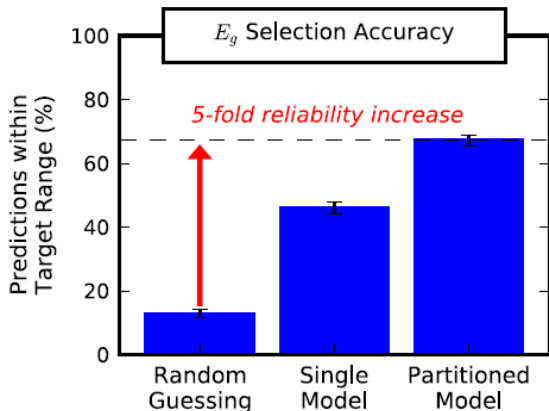


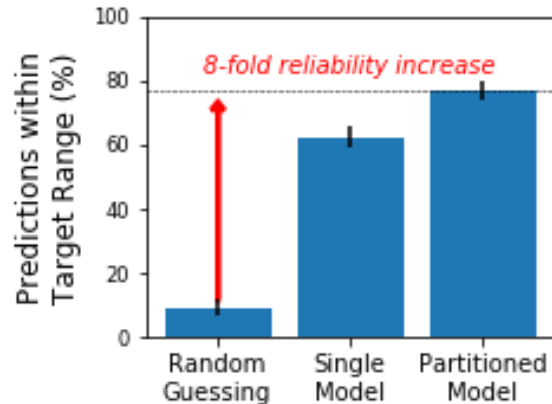
Table 2. Compositions and predicted band gap energies of materials predicted using machine learning to be candidates for solar cell applications

Composition	E_g (eV)
ScHg ₄ Cl ₇	1.26
V ₂ Hg ₃ Cl ₇	1.16
Mn ₆ CCl ₈	1.28
Hf ₄ S ₁₁ Cl ₂	1.11
VCu ₅ Cl ₉	1.19



Abbreviations: DFT, density functional theory; OQMD, open quantum materials database.

Replication in 2017



Out[11]:

	Entry	bandgap_predicted
1037	CoB2F9	1.380256
3414	YbAs7Cl6	1.156973
3884	Tl3OsO3.5	1.078918
1920	Cs8CoSe5	1.074358
421	Mg6SiTe8	1.119002

Better reproducibility via publication

What do I need to publish?

- *Datasets*: In a well-described format
- *Scripts*: Not just the core methods
- *Outputs*: Exact version from the paper
- *Models*: In a user-friendly way



Datasets: Citrination

- Or another database that serves *structured data*

Bulk Metallic Glasses

ID: 156839 - Version 1 - [Create new version](#) - [Edit](#) - [Delete](#)

Description:

Dataset associated with "A Machine Learning Approach for Engineering Bulk Metallic Glass Alloys." Contains the glass-forming ability (either bulk, ribbon, or none), critical casting diameter, supercooled liquid range, and glass transition temperature for many metallic alloys.

Data Views containing this dataset:

This dataset has not been used in any views.

[Show Less](#)

Search this dataset

Material Name or Chemical Formula



Property Name

Units

[Advanced Search Options](#)

[Search](#)

Showing results 1 to 24 of 7093

[First](#) [Next](#)

$B_{12}Fe_{78}Mn_{10}$



Chemical formula: $B_{12}Fe_{78}Mn_{10}$
Glass forming ability: None

$B_{17}Mo_{14.5}Ni_{68.5}$



Chemical formula: $B_{17}Mo_{14.5}Ni_{68.5}$
Glass forming ability: Ribbon

Datasets: Citrination

Ti₃₄Zr₁₁Cu₄₇Ni₈

+ Add Data Dataset ↓ Down



Chemical Formula: Ti₃₄Zr₁₁Cu₄₇Ni₈

Properties

Glass forming ability: BMG

Data Type

EXPERIMENTAL

References

1.

Url: <http://www.sciencedirect.com/science/article/pii/S0925838808012206>

dT_x: 28.8 K

Data Type

EXPERIMENTAL

References

1.

Url: <http://www.sciencedirect.com/science/article/pii/S0925838808012206>

D_{max}: 4.5 mm

Data Type

EXPERIMENTAL

References

1.

Url: <http://www.sciencedirect.com/science/article/pii/S0925838808012206>


Data and metadata stored together






In a well-documented format



Data Is Easy to Use Outside of My Purposes

```
{
  "category": "system.chemical",
  "uid": "00078D0F27AEB092DF423435AB653A96",
  "properties": [
    {
      "name": "Glass forming ability",
      "scalars": [
        {
          "value": "BMG"
        }
      ]
    }
  ],
  "dataType": "EXPERIMENTAL",
  "references": [
    {
      "url": "http://www.sciencedirect.com/science/"
    }
  ]
}
```

Scripts GitHub and WholeTale

 WardLT Merge branch 'master' of github.com:fang-ren/Discover_MG_CoVZr Latest commit a6882e9 on Apr 12

..		
 README.md	Documentation updates, updating results	8 months ago
 compare-models.ipynb	Updated results	8 months ago
 make-model.in	Added missing script	8 months ago
 run-all.bs	Autodetect number of processors	6 months ago
 run-cv-test.scala	Test whether new data changes best method for adding processing	11 months ago

 README.md 

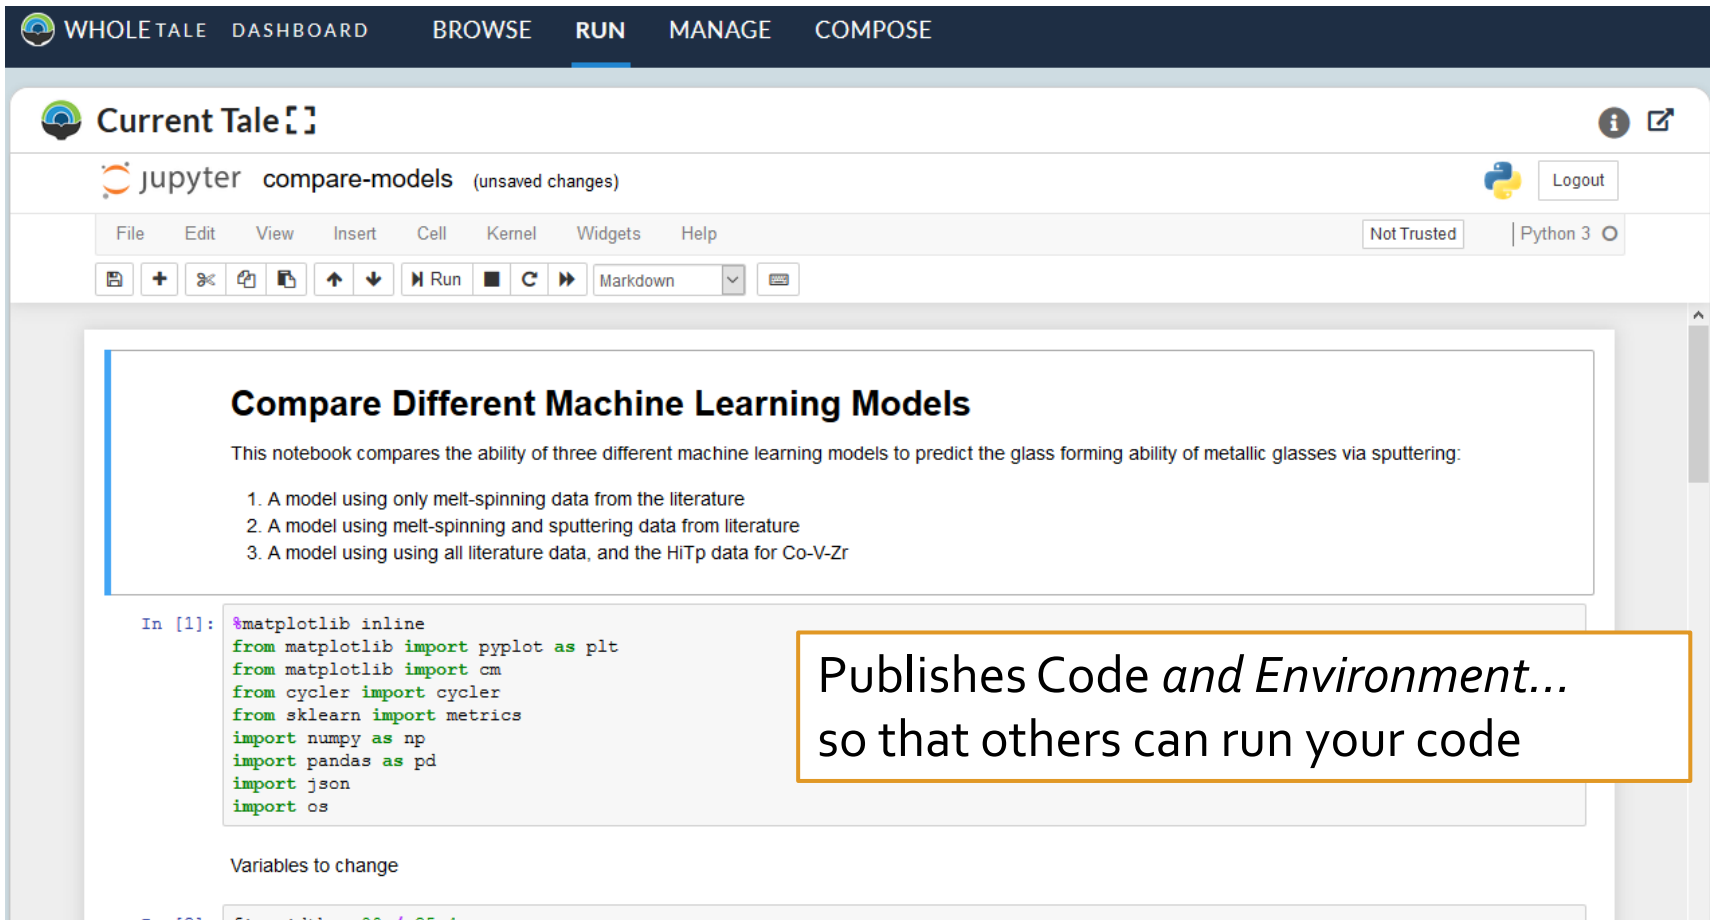
This directory contains a script designed to test how adding more data changes the accuracy of our machine learning model, and a script for running the model to find new metallic glasses.

To run the scripts call `./run_all.bs`.

The code for comparing the performance of the model as more data is contained within the notebook, `compare-models.ipynb`. The new predicted glasses are summarized in `new-glasses_P0.95_dist0.10.csv`.

Logan Ward 15 February 2018 https://github.com/fang-ren/Discover_MG_CoVZr/

Scripts: WholeTale



The screenshot displays the WholeTale web interface. At the top is a dark navigation bar with links: WHOLE TALE, DASHBOARD, BROWSE, RUN, MANAGE, and COMPOSE. Below this is a header for the 'Current Tale' section, which includes a Jupyter logo, the notebook title 'compare-models' with '(unsaved changes)', a Python logo, and a 'Logout' button. A menu bar contains options: File, Edit, View, Insert, Cell, Kernel, Widgets, and Help. To the right of the menu is a 'Not Trusted' warning and a 'Python 3' selector. Below the menu is a toolbar with icons for saving, adding cells, undo, redo, running, and other functions. The main content area shows a Jupyter notebook with the title 'Compare Different Machine Learning Models'. The notebook's text describes the goal of comparing three machine learning models for predicting the glass forming ability of metallic glasses via sputtering. It lists three models: 1. A model using only melt-spinning data from the literature; 2. A model using melt-spinning and sputtering data from literature; 3. A model using all literature data, and the HiTp data for Co-V-Zr. Below the text is a code cell with the following Python code:

```
In [1]: %matplotlib inline
from matplotlib import pyplot as plt
from matplotlib import cm
from cycler import cycler
from sklearn import metrics
import numpy as np
import pandas as pd
import json
import os
```

Below the code cell, it says 'Variables to change'. A text box on the right side of the notebook content area states: 'Publishes Code *and Environment...* so that others can run your code'.

Data: Materials Data Facility

If you use this dataset please add this citation to your publication:

Fang, Ren; Ward, Logan; Williams, Travis; Laws, Kevin J.; Wolverton, Christopher; Hatrick-Simpers, Jason; Mehta, Apurva, "Accelerated Discovery of Metallic Glasses through Iteration of Machine Learning and High-Throughput Experiments," 2018, <http://dx.doi.org/doi:10.18126/M2B06M>



MATERIALS
DATA
FACILITY

NIST

CHMaD

Title: Accelerated Discovery of Metallic Glasses through Iteration of Machine Learning and High-Throughput Experiments

Authors: [Fang, Ren](#)
[Ward, Logan](#)
[Williams, Travis](#)
[Laws, Kevin J.](#)
[Wolverton, Christopher](#)
[Hatrick-Simpers, Jason](#)
[Mehta, Apurva](#)

Issue Date: 16-Feb-2018

Publisher: Materials Data Facility

URI: <http://dx.doi.org/doi:10.18126/M2B06M>

Appears in Collections: [MDF Open](#)

Endpoint and path to dataset

82f1b5c6-6e9b-11e5-ba47-22000b92c6ec/published/publication_992/

Show full record

[Return to data publication dashboard](#)



Endpoint



Path

Go

select all

up one folder

refresh list

permissions

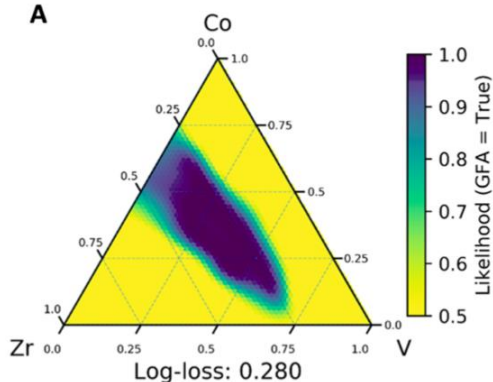
	plots	Folder
	results	Folder
	README.md	461 B
	all-training-data.obj	30.65 MB
	compare-models.ipynb	55 KB
	gfa-data.obj	22.32 MB
	gfa-model.obj	1.40 MB
	gfa-training-data.obj	30.65 MB
	make-model.in	2.90 KB
	make-model.out	3.96 KB
	new-glasses.json.gz	27.51 MB
	new-glasses_P0.95_dist0.10.csv	111.89 KB
	prediction-analysis.txt	317 B
	run-HiTp-data.out	1.97 KB
	run-all.bs	985 B
	run-cv-test.out	1.56 KB
	run-cv-test.scale	4.31 KB

Models: DLHub

Accelerated discovery of metallic glasses through iteration of machine learning and high-throughput experiments

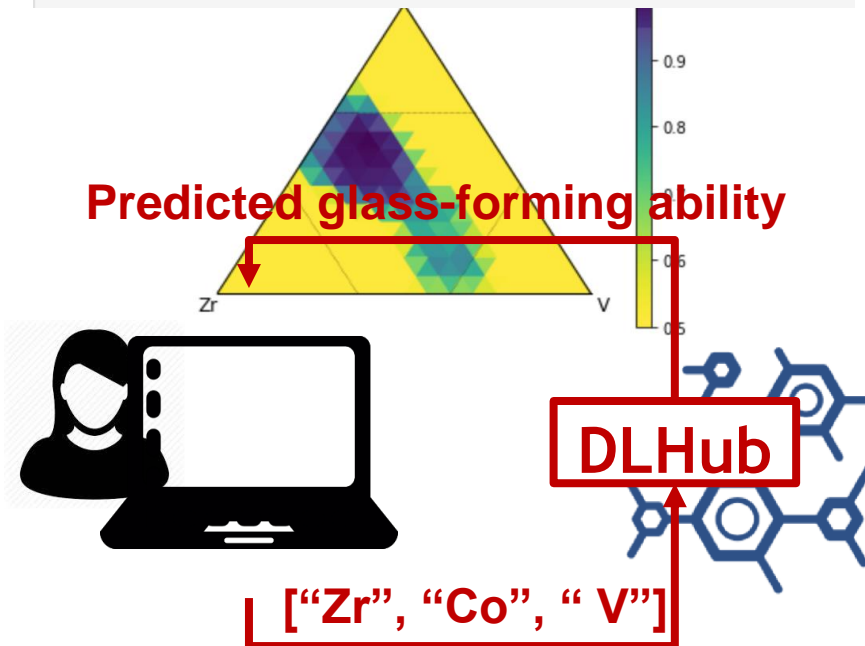
Fang Ren^{1,*}, Logan Ward^{2,3,*}, Travis Williams⁴, Kevin J. Laws⁵, Christopher Wolverton², Jason Hattrick-Simpers⁶ and Apurva Mehta^{1,†}

10.1126/sciadv.aag1566



```
servable_name = "metallic_glass"  
servable_id = dl.get_id_by_name(servable_name)  
elems = ["V", "Co", "Zr"]  
  
res = dl.run(servable_id, {"data": elems})
```

Predicted glass-forming ability



Overview of Today

1. Write notebooks like papers
2. Publish everything with your papers

Advice:

- Work in Jupyter until you're comfortable
- Read/watch "[I don't like Jupyter notebooks](#)"
- Learn [Markdown](#)

Simple Test: Channel

Here, we have a particle traveling forward or backwards along the channel of FCC AI. The path forward and backwards are identical, so we should get the same stopping power.

```
In [5]: %ipynb('python', dfk)
def compute_stopping_power(starting_point, direction, traj_computer=traj_computer):
    return traj_computer.compute_stopping_power(starting_point, direction, 1)
```

Set up calculations for the stopping power in the channel

```
In [6]: forward = compute_stopping_power([0,0,75,0,75], [1,0,0])
```

```
In [7]: backward = compute_stopping_power([0,0,75,0,75], [-1,0,0])
```

Wait for them to finish

```
In [8]: %time
forward = forward.result(); backward = backward.result()
```

CPU times: user 12 ms, sys: 12 ms, total: 24 ms
Wall time: 1.8 s

```
In [9]: print('Forward stopping power: ', forward[0])
print('Backward stopping power: ', backward[0])
print('Difference: ', backward[0]-forward[0])
```

Forward stopping power: 0.2343000200230004
Backward stopping power: 0.23431700358559353
Difference: 1.756276190512967e-05

Finding: Consistent with my initial expectations, they are indeed the same (within numerical tolerances).

